A Experimental setup

partially observed clinical simulator · GRPO with three reward formulations · one common held-out verifier

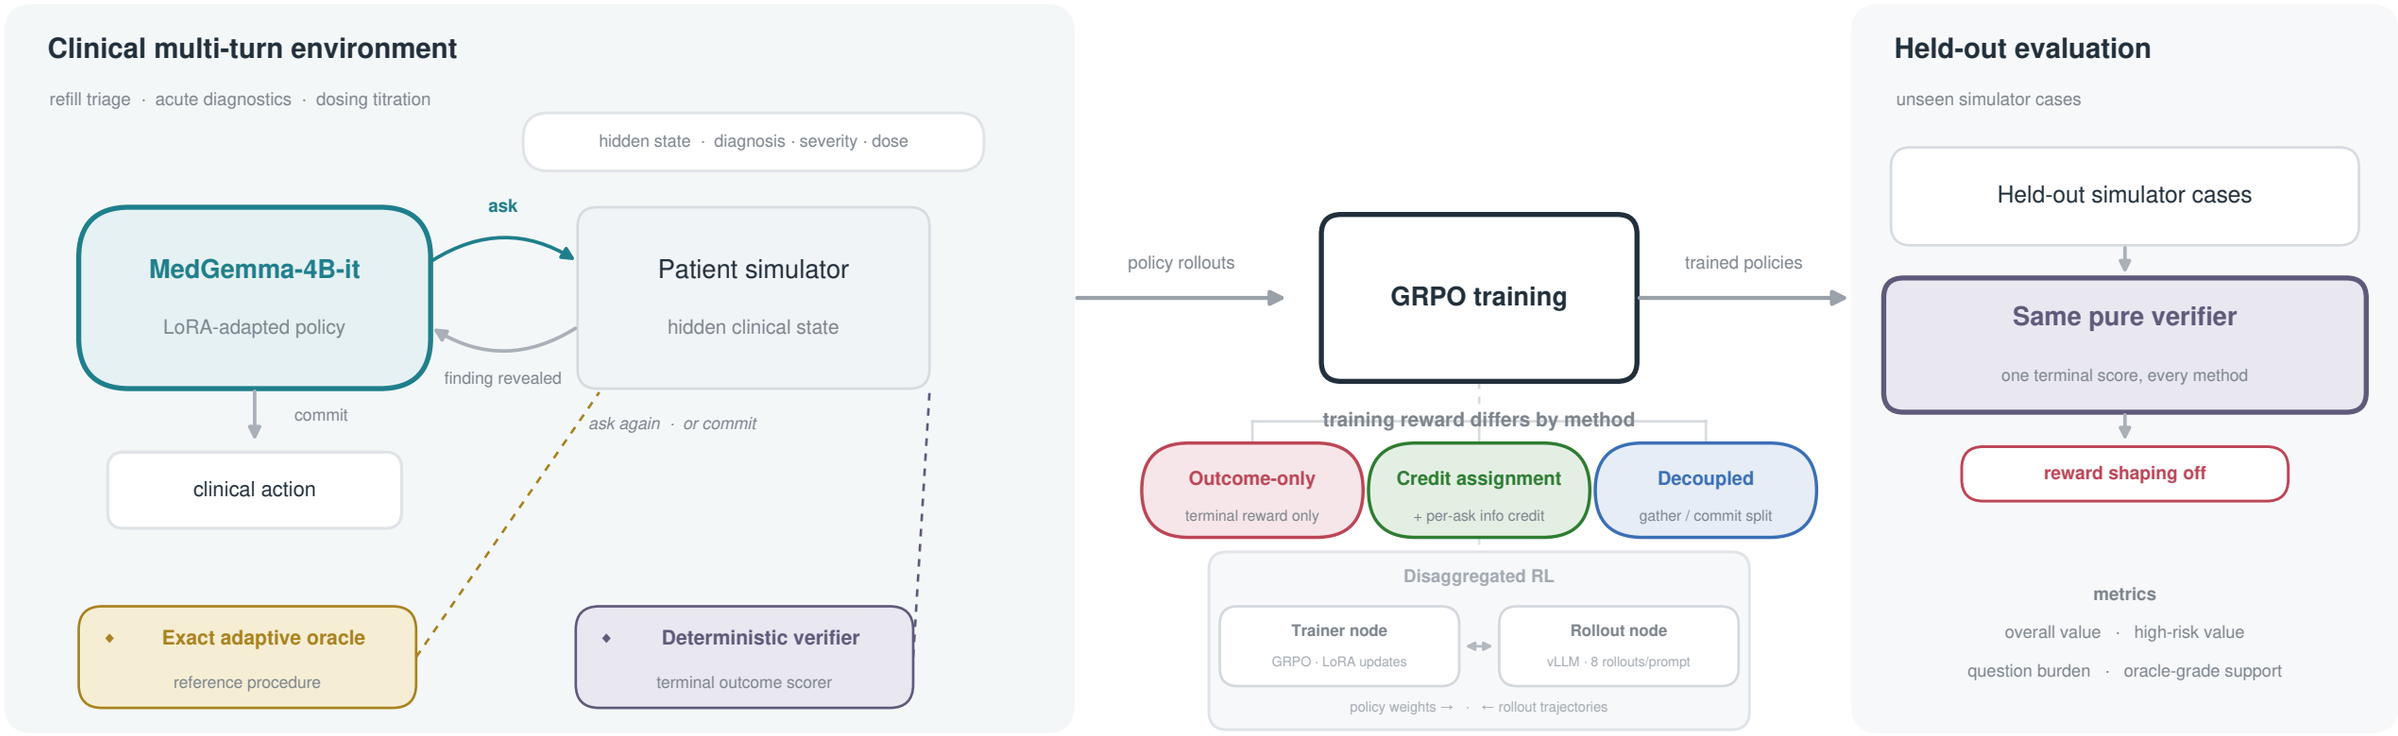

B Same patient, different decision procedures

one held-out high-risk case (acute abdominal pain, hidden emergent severity) · same simulator, same final verifier · only the training reward differs

single representative rollouts (k = 12 per condition); illustrative of procedure class, not modal or aggregate behaviour

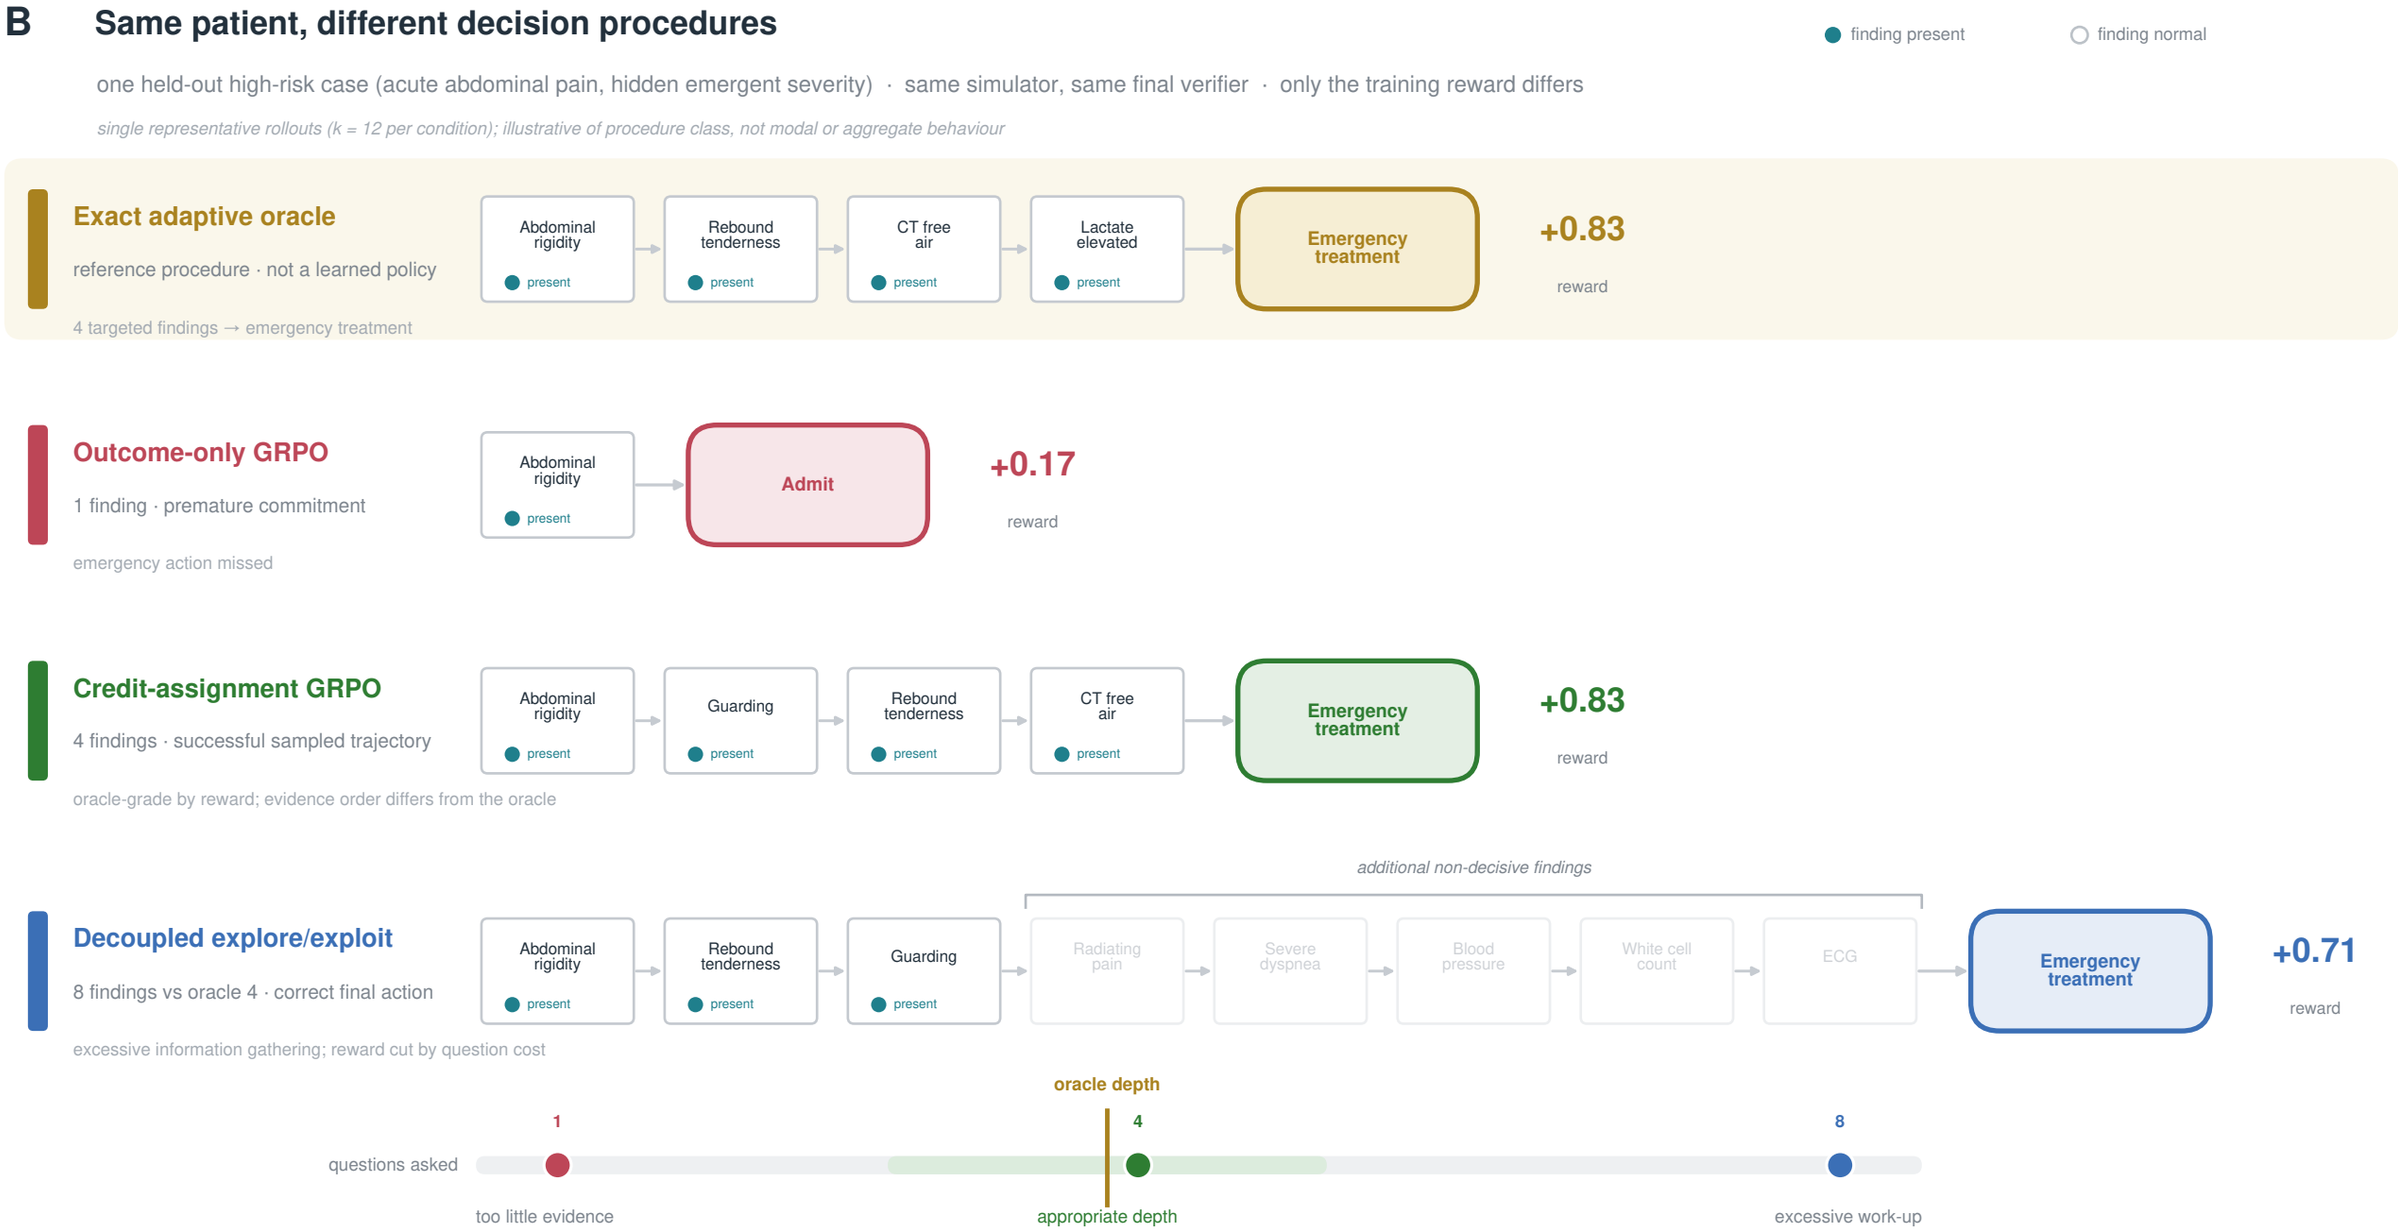

C Quantitative diagnostics results

held-out diagnostics · same unshaped terminal verifier for every method

High-risk value = mean verifier reward on the rare emergent subset (a behavioural measure, not a direct measure of clinical safety).

† single-seed reference (hatched). Outcome-only / Credit assignment: mean ± SD across 3 training seeds. Dashed line: exact adaptive oracle on the same held-out cases.